

Analysis of a Linear Regression Model for Time Series Data

D'Souza, Ajay
ajaydsouza@gatech.edu

Analysis of the suitability of a linear regression model to predict the log stock returns of one stock(GM) from the log stock returns of other stocks in the same Industry(Toyota,Ford)

Abstract

From the sample time series data set of daily log returns for GM, Toyota and Ford develop a Linear Regression Model to predict the returns for GM from the returns of Toyota and Ford. The log returns of GM is the response variable which the model should predict using the predictor variables, the log returns of Ford and Toyota.

Contents

List of Figures	3
List of Tables	3
1 Summary	4
2 Introduction	4
3 Body	5
3.1 Data Analysis	5
3.2 Analysis of the simple Linear Regression Model	8
3.2.1 Model Selection - Predictor Subset	8
3.2.2 Model Selection - AIC	10
3.2.3 Model Selection -ANOVA	10
3.2.4 VIF	11
3.2.5 Partial Residual Plot	11
3.2.6 Diagnostics - Leverage	13
3.2.7 Diagnostics - Influence	15
3.2.8 Diagnostics - Studentized Residuals	17
3.2.9 Verify Model Assumptions	18
3.2.10 Model Summary	20
3.3 Considering a Linear Model with Polynomial Terms	20
3.3.1 Model Selection - Predictor Subset	21
3.3.2 Model Selection - AIC	24
3.3.3 ANOVA	24
3.3.4 Diagnostics - Leverage	25
3.3.5 Diagnostics - Studentized Residuals	25
3.3.6 Model Assumptions	27
3.3.7 Compare Model Summary	29
3.4 10 Fold Cross Validation - Evaluate Models	29
4 Conclusion	32
Appendices	33
References	53

List of Figures

1	Correlation and Density Plots of of Toyota, Ford and GM log Returns	6
2	QQ Plots - of Toyota, Ford and GM log Returns Vs Normal and T Dis- tribution with $\nu = 4$	7
3	Plot of No of Predictor Variables Vs BIC, C_p , R^2	9
4	Partial Residual Plots for each predictor variable - Ford, Toyota	12
5	Plot of the Hat Values(Leverage) Vs Index	14
6	Square Root Cooks Distance (Influence) - Scatter Plot Vs Index and Half Normal Plots	16
7	Studentized Residuals Versus Leverage and Cooks Distance	17
8	Plots to check assumptions of Normality, Linearity and Constant Variance - Plot of Studentized Residuals Vs Predictors, Fitted Values	19
9	Poly Term Linear Model - Plot of No of Predictor Variables Vs BIC, C_p , R^2	23
10	Poly Term Linear Model - Studentized Residuals Versus Leverage and Cooks Distance	26
11	Poly Term Linear Model - Plots to check assumptions of Normality, Linear- ity and Constant Variance - Plot of Studentized Residuals Vs Predictors, Fitted Values	28
12	10 Fold Cross Validation Results - Comparison of Poly(Toyota,2) and Non Poly Term Models - Validation Set MSE for each Fold	31

List of Tables

1 Summary

We analyze the provided daily log returns sample data to understand factors like the relationship between the variables, the distributions of the variables. We fit a simple linear model with the first degree terms of two the predictor variables Ford and Toyota. The performance and assumptions of this model are analyzed. Based on this diagnosis we fit a second linear model with a additional second degree orthonormal Polynomial term on the log returns of Toyota. We analyze the performance of the two models. We compare the performance of the two models with 10-Fold cross validation. We use this analysis and validation set results to choose the better linear model.

2 Introduction

We are provided with a sample size of $N = 709$ data points. The Preliminary data analysis shows a strong correlation between the Ford and GM log returns. The correlation between Toyota and GM and Toyota and Ford while being positive is weak. The simple linear model fitted to this data has a $R^2 \approx 0.37$. ANOVA results show that almost all of the contribution to explaining the variation from the regression sum of squares in this model comes from the Ford predictor variable. For the predictor variable Toyota, both the F-value from ANOVA and the P-Value from the summary of the linear model are at around 11%. So the support for rejecting the NULL hypothesis and accepting that log returns of Toyota as a predictor are significant to this linear model is not very strong. We check if the linear regression model can be enhanced by a linear relationship with a higher order polynomial term in Toyota. We generate orthonormal polynomial terms up to degree 2 for Toyota and try to fit a linear regression model to this data. We check if the linear model fitted has a better R^2 . We also check the F-Value from ANOVA and the P-Value from the summary of the linear regression to see if it offers better support for rejecting the NULL hypothesis and accepting the 2nd degree polynomial terms for Toyota as significant to this model.

For both models we check for deviation from our model assumptions of normality, constant variance, linearity by using plots of the studentized residuals against fitted values, predictors, density and QQ plots with normal distribution. We check for high leverage points, residual outliers and data points with high influence using the plots for hat matrix values, Studentized Residuals and Cooks Distance.

Finally we do robust comparison of the 2 models using a 10 Fold cross validation. The Mean Residual Square Error (MSE) on the all the hold out validation test sets in each fold is used as the metric for evaluating model performance.

3 Body

3.1 Data Analysis

The sample dataset has $N = 709$ data points. The following is the correlation matrix for the log returns of Ford, Toyota and GM. There is a strong linear correlation of 0.61 between Ford and GM. There is a weak linear correlation of 0.20 between Toyota and GM and weak correlation of 0.24 between Toyota and Ford.

```
Correlation:
> cor(cbind(Toyota,Ford,GM))
      Toyota Ford  GM
Toyota  1.00 0.24 0.20
Ford    0.24 1.00 0.61
GM      0.20 0.61 1.00
```

The following figure (1) is the density and scatter plot of the daily log returns sample data. The scatter plot shows the positive correlation between Ford and GM and weak correlation between Toyota and GM and Toyota and Ford. The Density plots for each of the three variables shows a density distribution with thick tails. The GM and Ford data shows a slight skew. The Toyota distribution appears to be fairly symmetric.

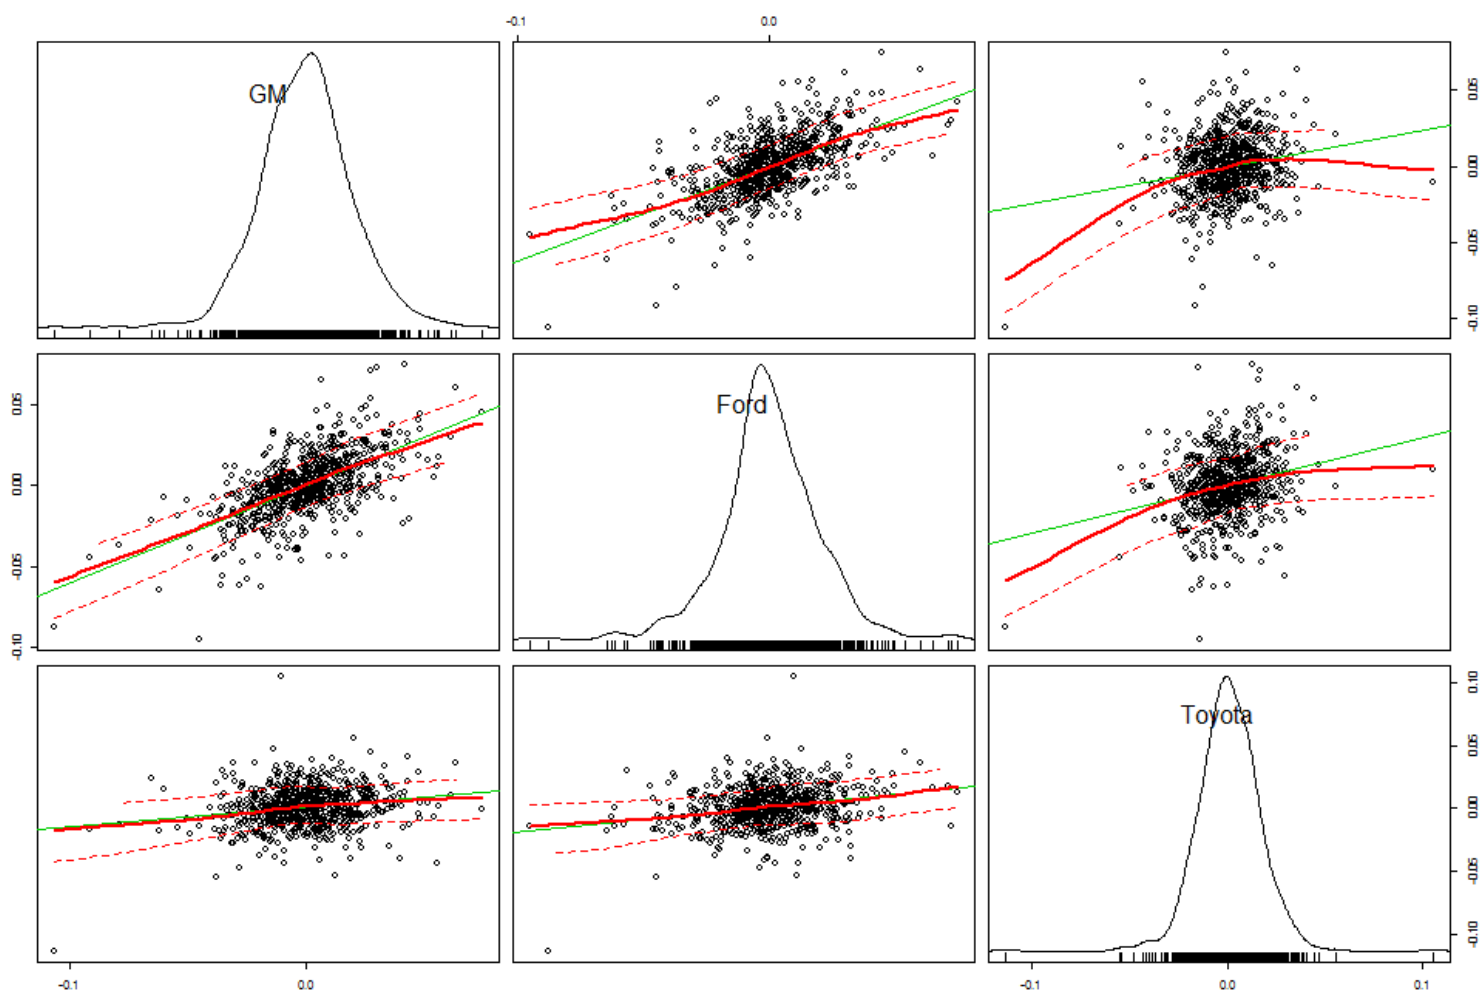


Figure 1: Correlation and Density Plots of of Toyota, Ford and GM log Returns

Figure (2) is the QQ plot for the sample data with both a Normal and T distribution with $\nu = 4$. All plots appear thick tailed. GM and Ford data seem to be closer in their distribution. The body and shoulders for Toyota plot is straighter and a better fit for both the normal and T distribution compared to GM and Ford data.

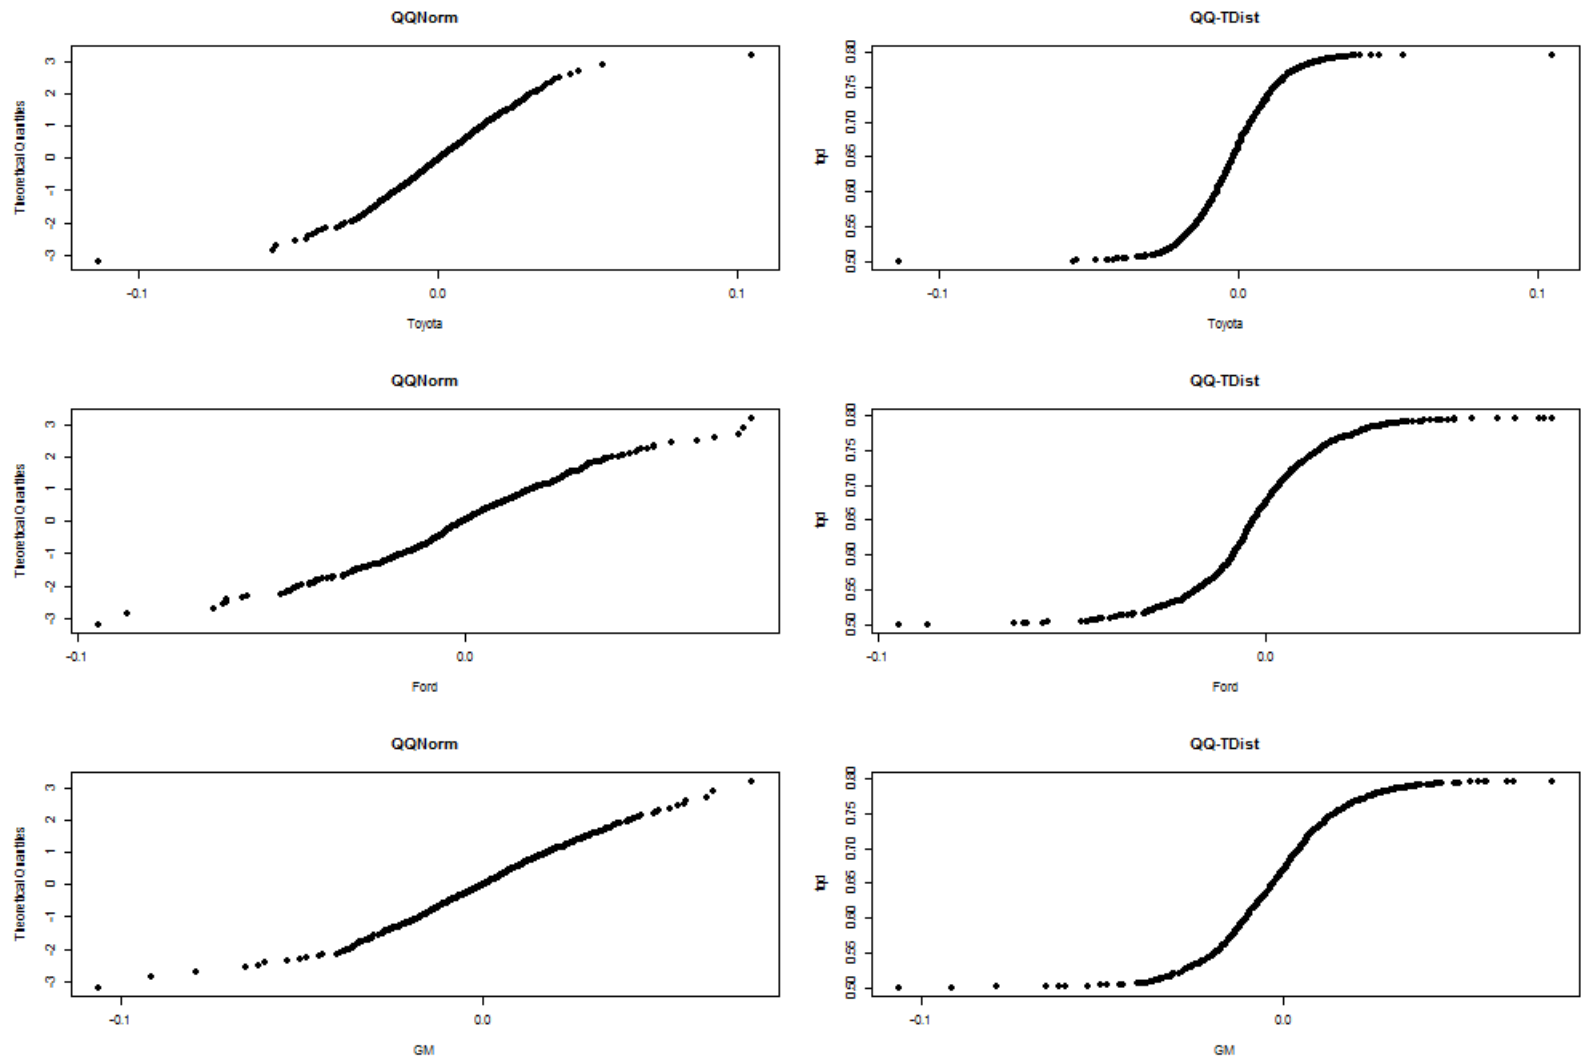


Figure 2: QQ Plots - of Toyota, Ford and GM log Returns Vs Normal and T Distribution with $\nu = 4$

3.2 Analysis of the simple Linear Regression Model

3.2.1 Model Selection - Predictor Subset

Since we have just 2 predictor variables, model selection for the simple linear regression model simply involves choosing either the best 1 predictor variable model or the 2 predictor model. The *leaps()* call in R, chooses the best k-subset of predictor variables for each $k \in 1 \dots 2$. The following are the results of *leaps()* function. The best k-Subset predictors are listed along with their C_p values.

Leaps - Best Model for each K Predictor by C_p :

	Ford	Toyota	
1	1	0	3.6
2	1	1	3.0

Similar selection of the best k-subset of predictors can be performed using the *regsubsets()* call in R. Results are attached below. In both cases as expected we see that for a 1 predictor model, Ford is the better linear predictor for GM log returns than Toyota.

Regression Subsets - Choosing the Best set of Predictors:

Subset selection object

Call: *regsubsets.formula*(data\$GM ~ ., data = as.data.frame(cbind(data\$Ford, data\$Toyota)), nbest = 1)

2 Variables (and intercept)

Forced in Forced out

V1 FALSE FALSE

V2 FALSE FALSE

1 subsets of each size up to 2

Selection Algorithm: exhaustive

V1 V2

1 (1) "*" " "

2 (1) "*" "*"

We choose the best subset of predictors, among these two 1-Predictor and 2-Predictor models by comparing the \mathbf{R}^2 results and along with AIC, BIC and \mathbb{C}_p model quality measures. For \mathbf{R}^2 , a higher value is better, for the AIC, BIC and \mathbb{C}_p , a lower value indicates a better model. Figure (3) plots these values.

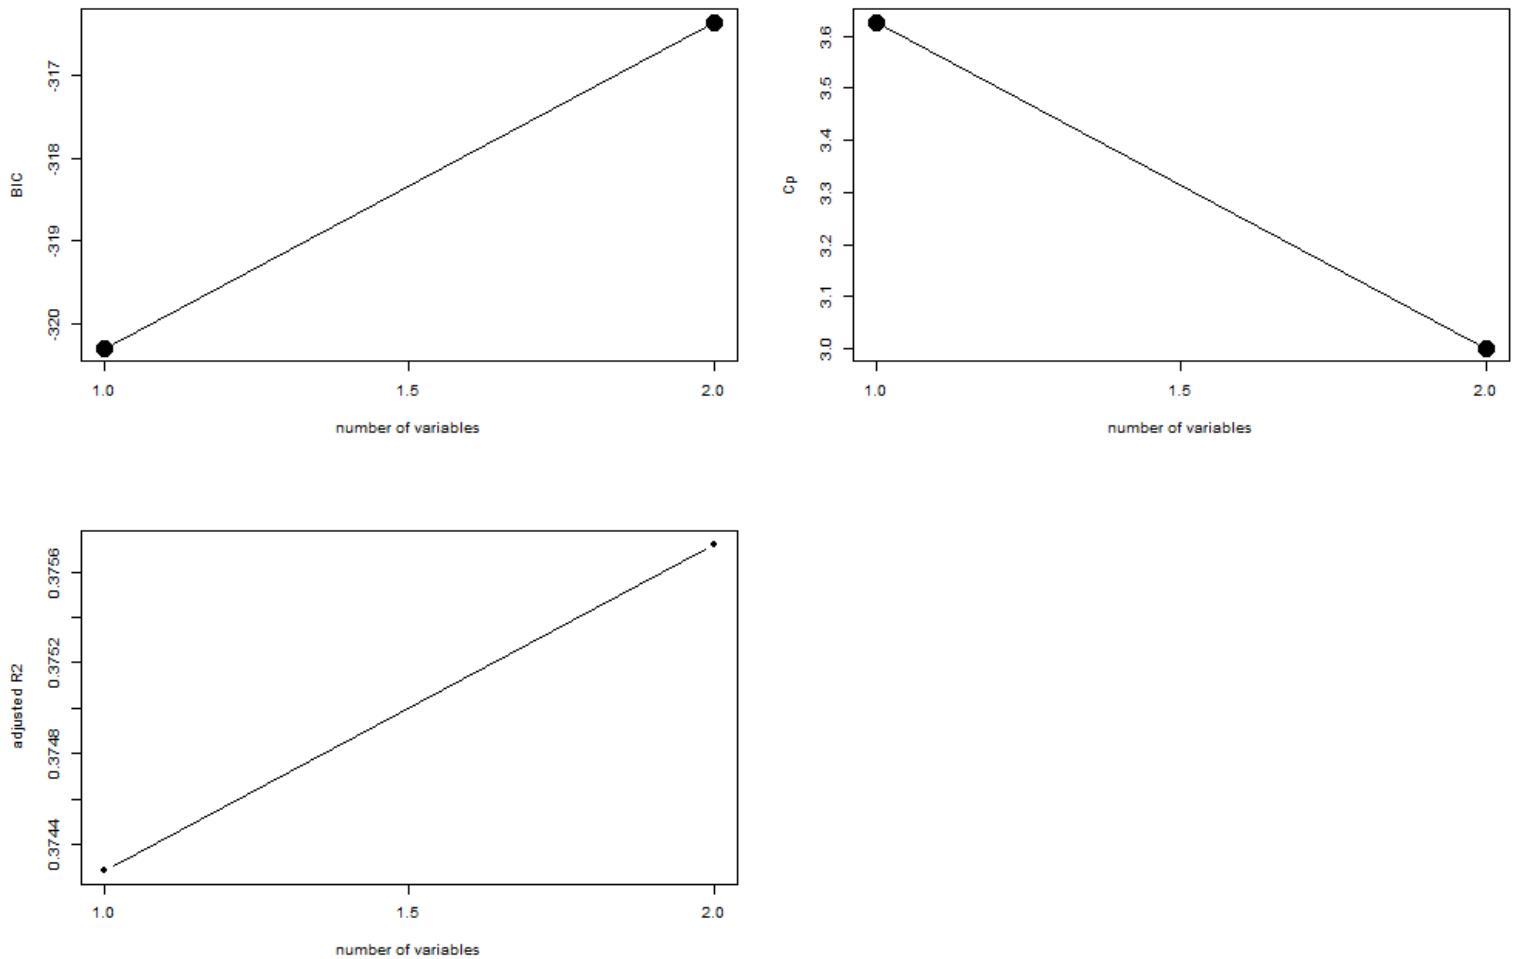


Figure 3: Plot of No of Predictor Variables Vs BIC, \mathbb{C}_p , \mathbf{R}^2

3.2.2 Model Selection - AIC

The most parsimonious model is the model with the lowest AIC value. The *stepAIC()* R call systematically removes one predictor variable at a time till the AIC stops reducing. The following results of the *stepAIC()* call show that the 2 predictor variable model has the lowest AIC of -5885.44.

```
Step AIC  GM ~ Ford + Toyota :  
Start:  AIC=-5885.44  
GM ~ Ford + Toyota
```

	Df	Sum of Sq	RSS	AIC
<none>			0.1745	-5885
- Toyota	1	0.00065	0.1752	-5885
- Ford	1	0.09515	0.2697	-5579

All of the model selection criteria considered above suggest that the model having both the predictors Ford and Toyota is a slightly better choice than the 1 predictor model with only Ford as the predictor.

3.2.3 Model Selection -ANOVA

We can use ANOVA to compare the model with only Ford as the predictor with the 2-predictor model having both Ford and Toyota. The ANOVA results show that the $R^2 \approx 0.37$. So close to 37% of the total variance in GM can be explained by the predictor model having both Ford and Toyota as predictors. The increase regression sum of squares with the 2-predictor model is 0.0006492. This is very small compared to the regression sum of squares of 0.10518 from the model with only Ford. The F-Value for the Null hypothesis is $0.106 \approx 11\%$. So we can accept the 2-predictor model as significant and reject the Null hypothesis at the 11% level, which is weak. A smaller F-value would have suggested a stronger support for the 2-predictor model.

```
ANOVA for the linear regression model : GM ~ Ford + Toyota :  
Analysis of Variance Table
```

```
Response: GM
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Ford	1	0.10518	0.10518	425.482	<2e-16 ***
Toyota	1	0.00065	0.00065	2.626	0.106

```
Residuals 706 0.17453 0.00025
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

3.2.4 VIF

The Variance Inflation Factor (VIF) for both Ford and Toyota are very close to 1. This confirms what we have seen earlier with the correlation that the predictor variables Ford and Toyota are not very well correlated to one another and one cannot be predicted from the other. This means that the standard deviation of the estimated slope of one is not much affected by the presence of the other in the model. So the quality of prediction of the response is not much affected by having both the predictor variables in the model.

Variance Inflation Factor:

	Ford	Toyota
	1.06246	1.06246

3.2.5 Partial Residual Plot

Figure (4) plots the partial residuals for his model. The plot in (b) has a slightly shallower slope than in plot (d). This shows that when predictor variable Ford is in the model, the effect of Toyota on the response variable GM is less than when Ford is not in the model. This is due to the slight collinearity that exists between the Ford and Toyota Predictor Variables. The same effect but to a much smaller degree can be seen in plots (a) and plot(c) for Ford and GM.

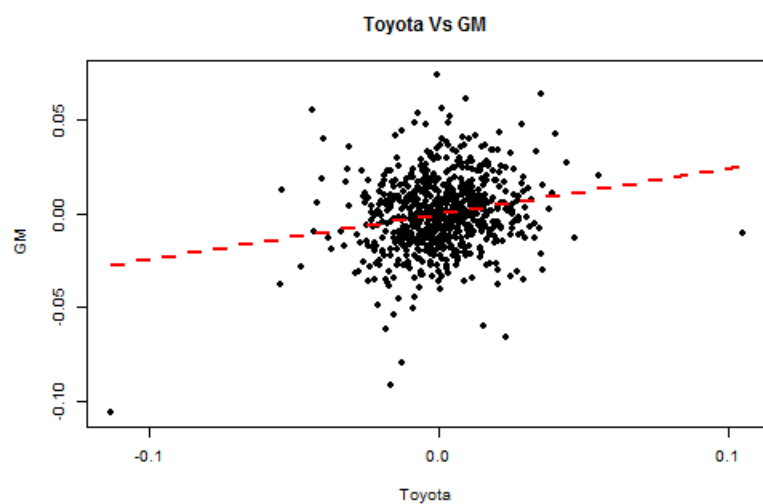
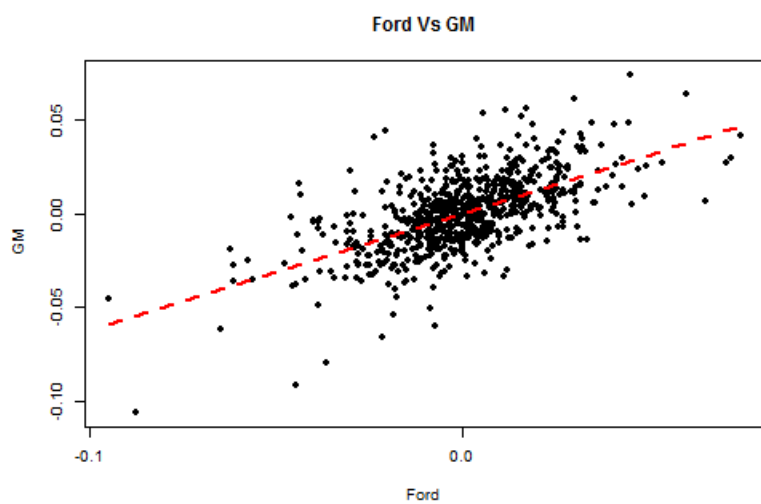
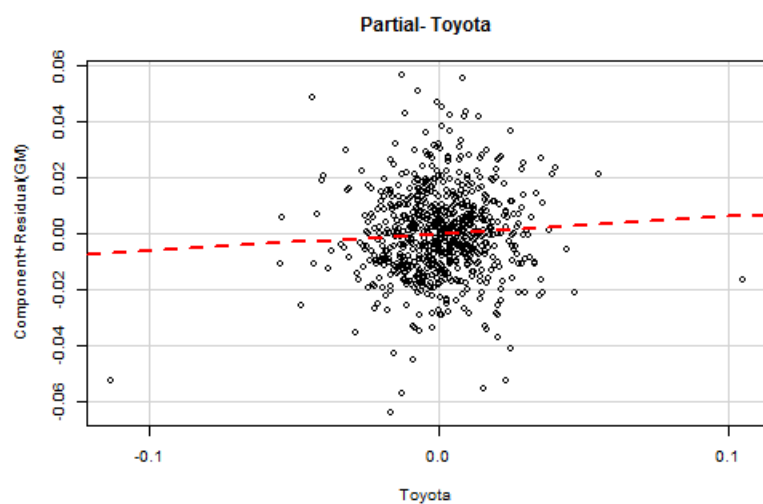



Figure 4: Partial Residual Plots for each predictor variable - Ford, Toyota

3.2.6 Diagnostics - Leverage

Leverage measures the influence a point has on its own fitted value. Figure (5) plots leverage for each of the 709 data points. There 59 high leverage points with leverage $\geq \frac{2(p+1)}{n} \approx 0.01$ where $\frac{(p+1)}{n}$ is the average leverage. These high leverage data points show a higher correlation between Toyota, Ford and GM than seen in the overall data. Since the predictors and the response have a thick tailed distributions, this indicates tail dependence between the returns when returns are extreme.

The Correlation between the 59 High Leverage Data Points:

	Toyota	Ford	GM
Toyota	1.00	0.25	0.26
Ford	0.25	1.00	0.77
GM	0.26	0.77	1.00

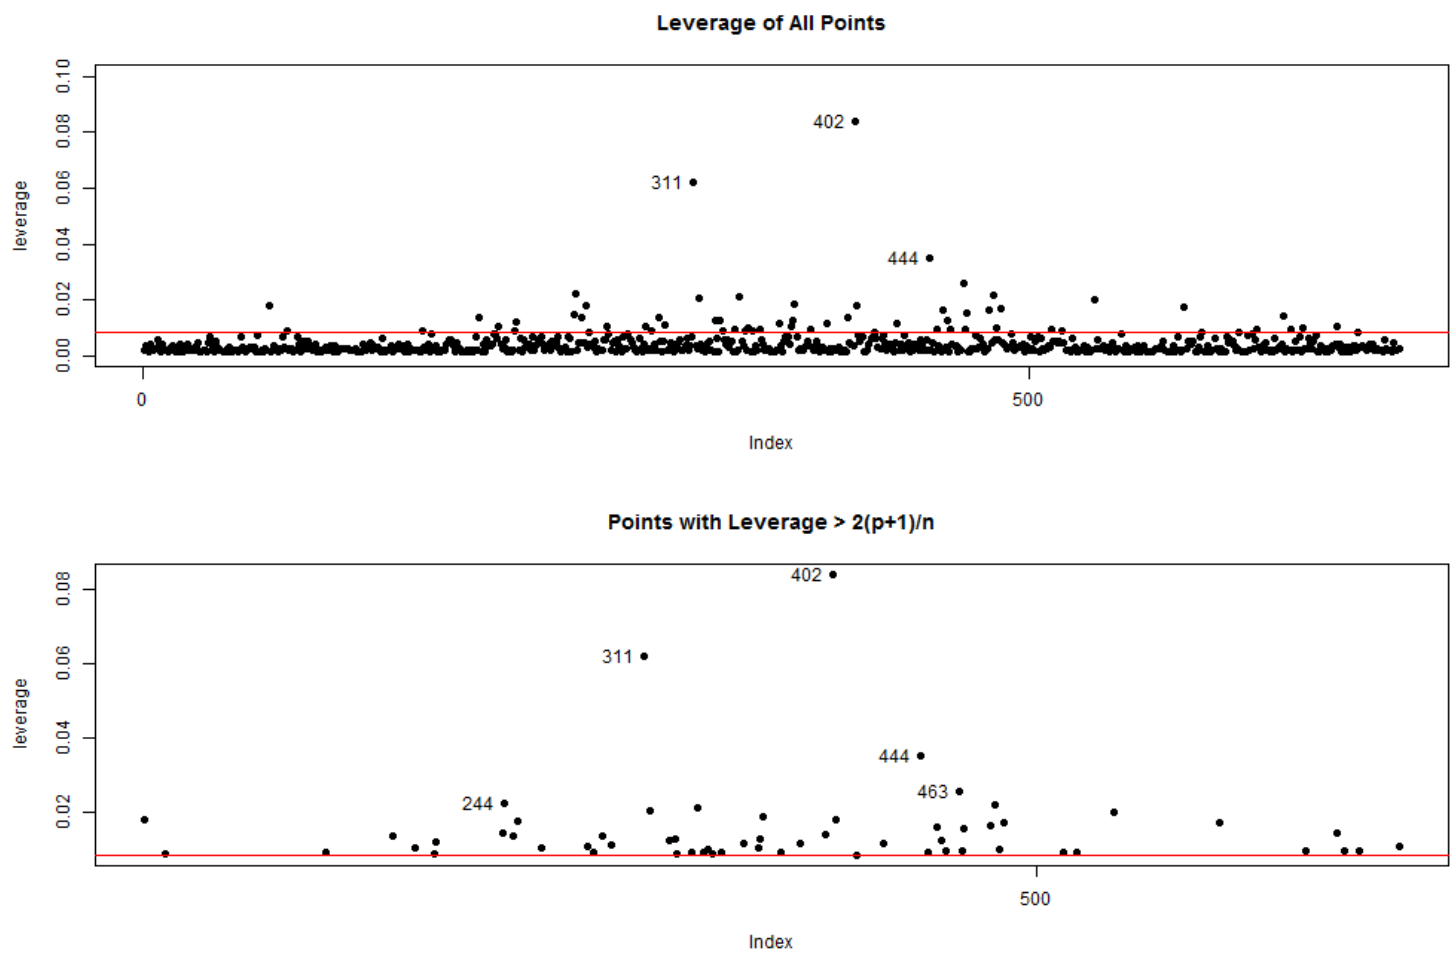


Figure 5: Plot of the Hat Values(Leverage) Vs Index

3.2.7 Diagnostics - Influence

Cook's distance measures the influence of a observation point on the fitted values. Figure (6), the scatter plot of the Square root of the cooks distance and the half normal plot show observation 402 to be having a largest influence. From the Leverage plot and the Studentized Residual plot we notice that this observation is also an residual outlier and a high leverage point. The following is the data for observation 402.

	Toyota	Ford	GM	leverage
402	-0.113	-0.0875	-0.106	0.493

The data shows large negative returns that seem to have affected all 3 stocks for this day. This seems to be a valid data point, showing tail dependence and that needs to be factored into the model.

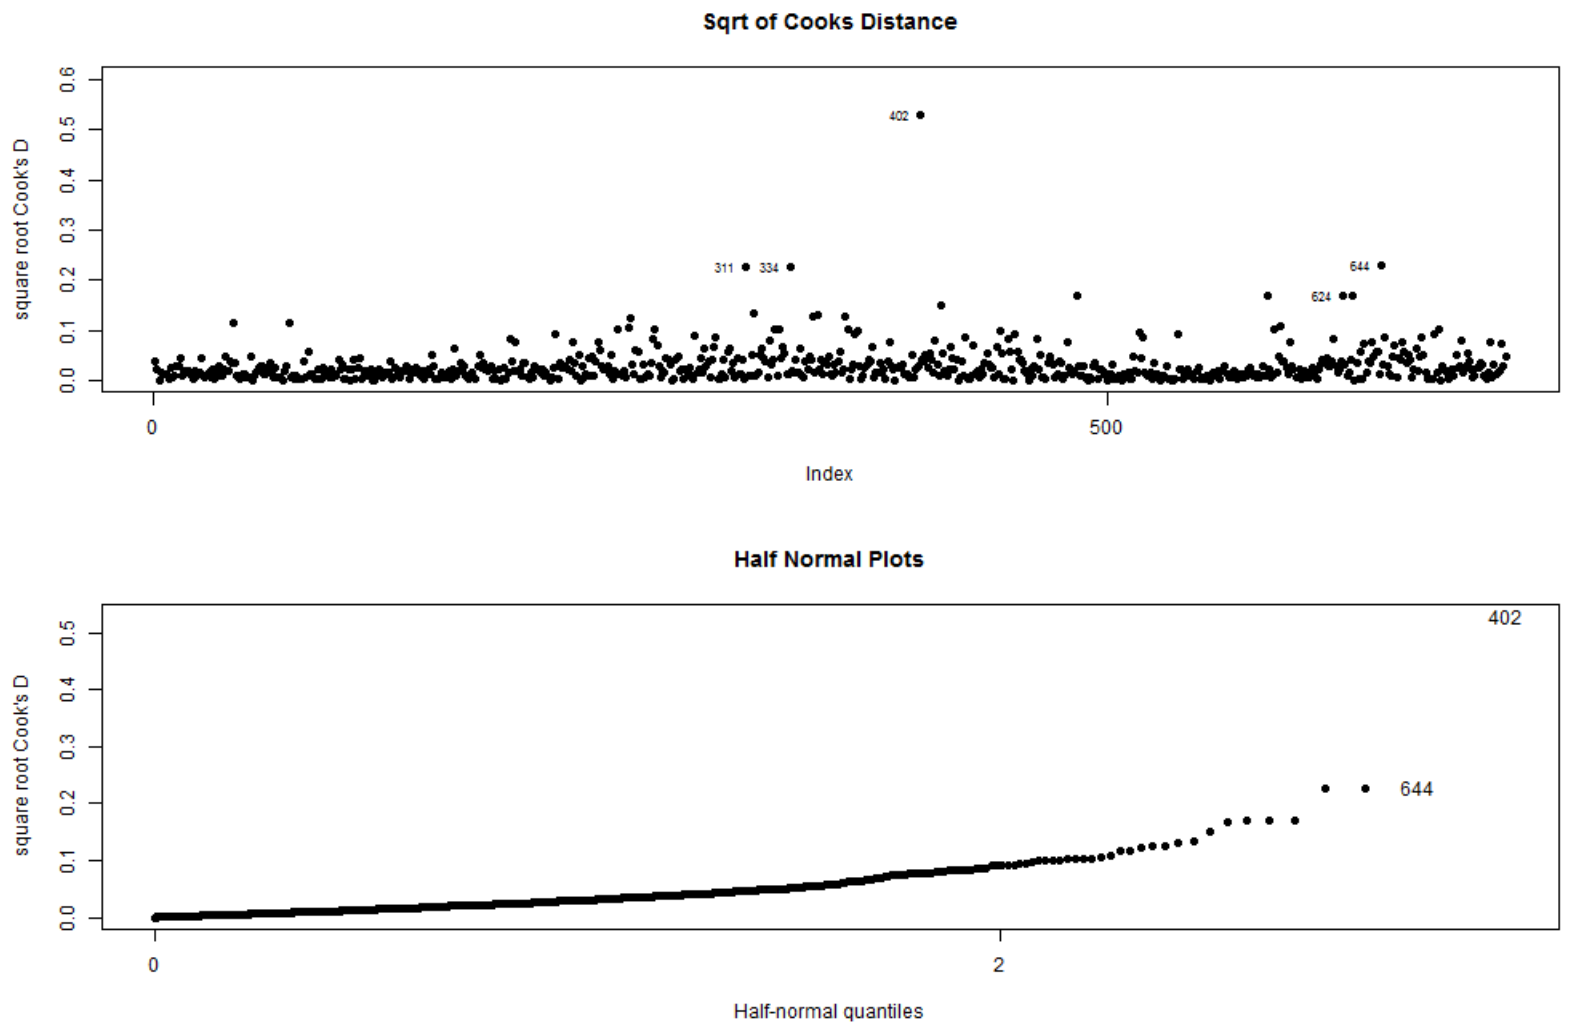


Figure 6: Square Root Cooks Distance (Influence) - Scatter Plot Vs Index and Half Normal Plots

3.2.8 Diagnostics - Studentized Residuals

Raw residuals do not have constant variance since they depend on the leverage and residual outlier position of that point. To overcome this we use Standardized Residuals also called externally Studentized Residuals. Figure (7) plots the Studentized Residuals for each point Versus the Leverage and Cooks Distance for that point. High Leverage points have lower residuals and the variation in $R_{Student}$ increases as the influence of the points increase.

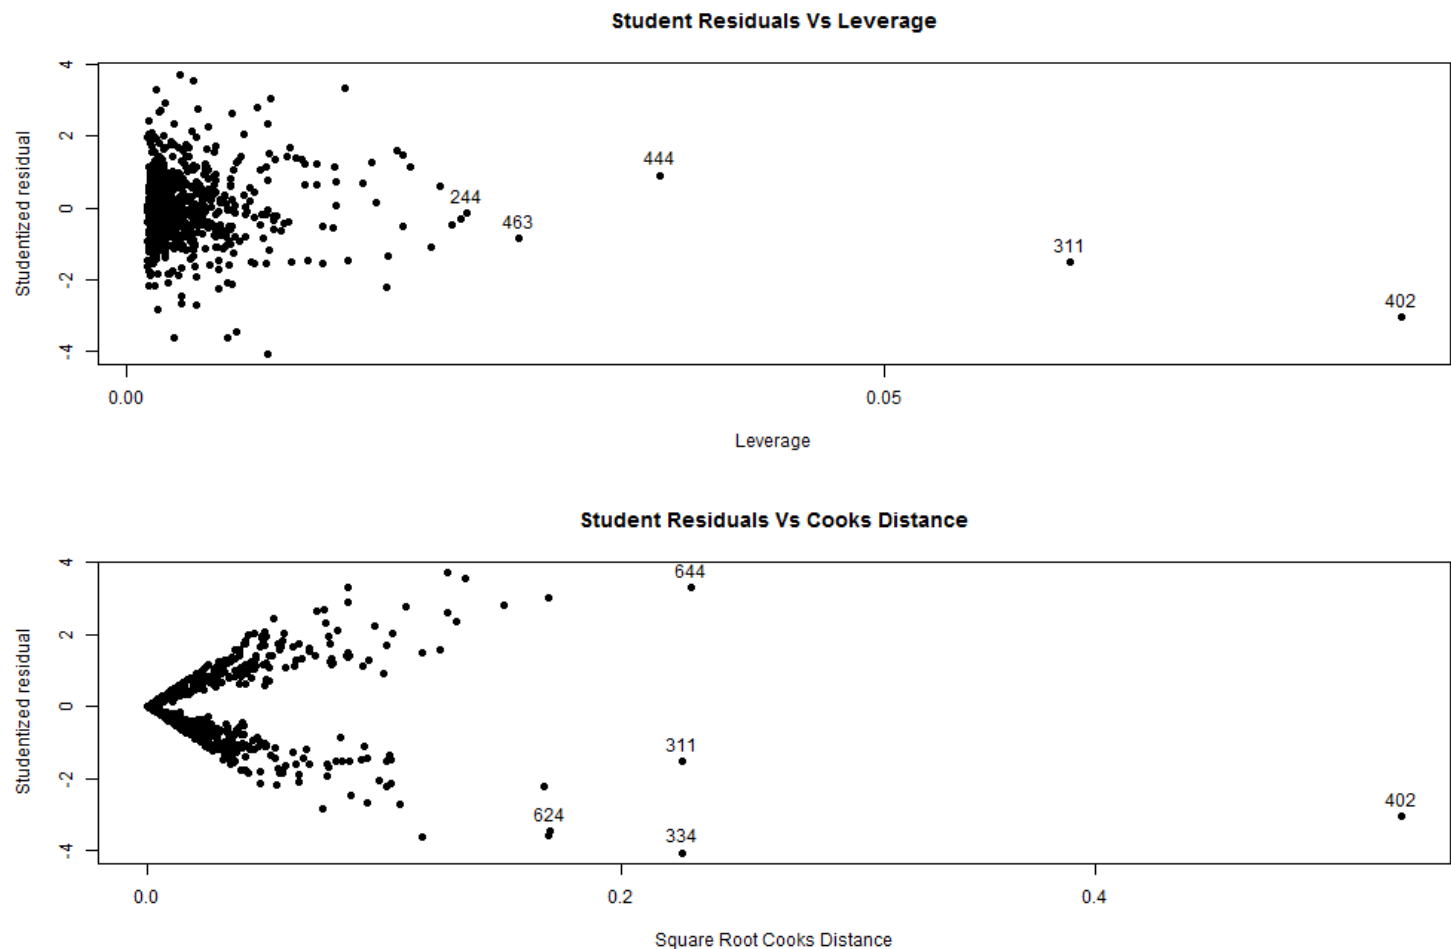


Figure 7: Studentized Residuals Versus Leverage and Cooks Distance

3.2.9 Verify Model Assumptions

Figure (8) plots the metrics that enable checking of the model assumptions of normality, linearity and constant variance. The density plot and the Normal QQ plot of the studentized residuals shows a thick tailed symmetric distribution. The plot of the Studentized Residuals versus Ford shows linearity, but the plot of the Studentized Residuals versus Toyota shows slight nonlinearity at the extremes. This would indicate tail dependence as discussed earlier at extreme values of returns. The plot of the Studentized Residuals versus Fitted Values is fairly linear and horizontal showing constant variance.

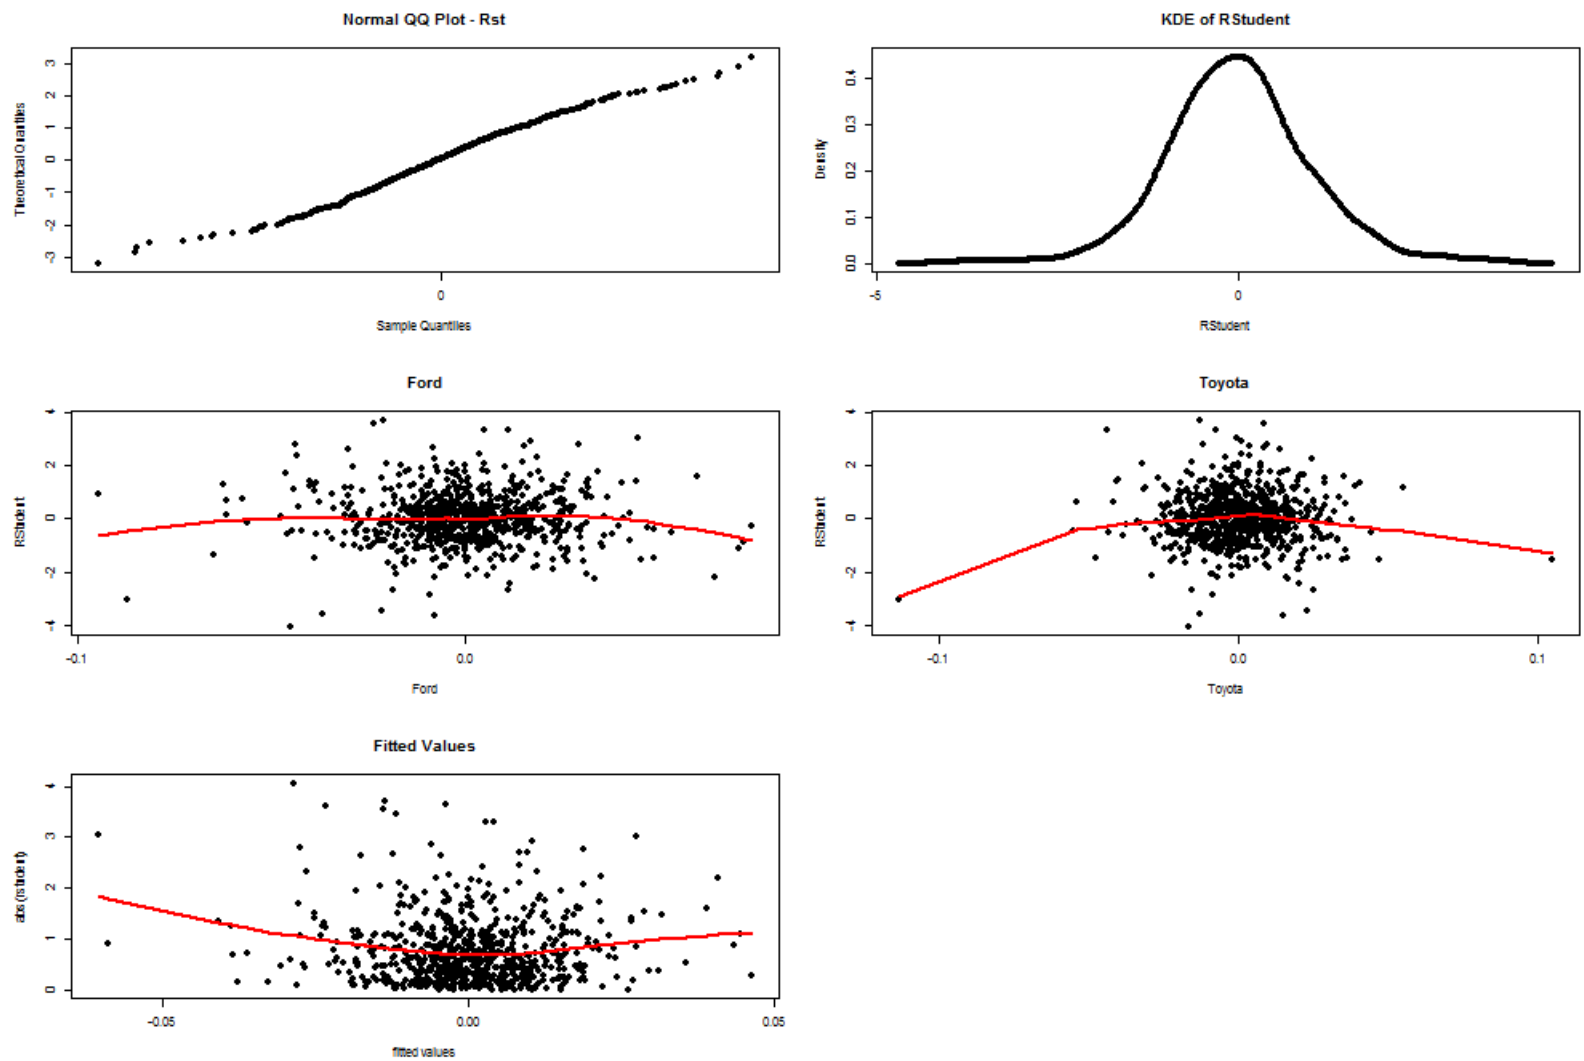


Figure 8: Plots to check assumptions of Normality, Linearity and Constant Variance - Plot of Studentized Residuals Vs Predictors, Fitted Values

3.2.10 Model Summary

The following is the summary of the simple linear regression model with 2-Predictors variables Ford and Toyota. The $R^2 \cong 0.38$ for this model.

Summary for the linear regression model : GM ~ Ford + Toyota :

Call:

```
lm(formula = GM ~ Ford + Toyota, data = data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.06285	-0.00965	-0.00040	0.00898	0.05752

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	7.05e-05	5.91e-04	0.12	0.91
Ford	6.14e-01	3.13e-02	19.62	<2e-16 ***
Toyota	6.13e-02	3.78e-02	1.62	0.11

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.0157 on 706 degrees of freedom

Multiple R-squared: 0.377, Adjusted R-squared: 0.376

F-statistic: 214 on 2 and 706 DF, p-value: <2e-16

3.3 Considering a Linear Model with Polynomial Terms

We have seen that the correlation between Toyota and GM is very low at 0.20. The Studentized Residual Plot versus the predictor variable Toyota in section (3.2.9) shows a slight deviation from linearity at the extremes. One way to overcome this non linearity in the Studentized Residuals with predictor variable Toyota is to see if we could fit a better linear model using degree 2 polynomial terms in Toyota. The polynomial terms used will be generated by the *poly()* function in R and will be orthonormal, as in QR distribution.

The following is the correlation matrix with degree 2 polynomial terms in Toyota. It can

be seen that the correlation between the 2nd degree term in Toyota and Ford is -0.075 , which is far less than the correlation of 0.24 between the degree 1 term in Toyota and Ford.

```
> cor(cbind(poly(Toyota,2),Ford,GM))
      1      2  Ford  GM
1  1.0e+00 -1.9e-17 0.242 0.20
2  -1.9e-17  1.0e+00 -0.075 -0.13
Ford  2.4e-01 -7.5e-02  1.000  0.61
GM    2.0e-01 -1.3e-01  0.613  1.00
```

3.3.1 Model Selection - Predictor Subset

The results of both the *leaps()* and *regsubsets()* functions in R which choose the best k-subset of predictors for each k are shown below. Results show that both functions choose the 2nd degree polynomial term in Toyota over the first degree term in Toyota for the 2-subset of predictor model.

Leaps - Best Model for each K Predictor by Cp:

	Ford	Toyota1	Toyota2	
1	1	0	0	11.2
2	1	0	1	4.8
3	1	1	1	4.0

Regression Subsets - Choosing the Best set of Predictors by BIC/Cp:

Subset selection object

Call: *regsubsets.formula*(data\$GM ~ ., data = as.data.frame(cbind(data\$Ford, poly(data\$Toyota, 2))), nbest = 1)

3 Variables (and intercept)

Forced in Forced out

V1 FALSE FALSE

'1' FALSE FALSE

'2' FALSE FALSE

1 subsets of each size up to 3

Selection Algorithm: exhaustive

	V1	'1'	'2'
1	(1)	"*" " " " "	
2	(1)	"*" " " "*"	
3	(1)	"*" "*" "*"	

We choose the best among these two 1,2 and 3-Predictor models by comparing the \mathbf{R}^2 results and along with AIC, BIC and \mathbb{C}_p model quality measures. For \mathbf{R}^2 a higher value is better, for the AIC, BIC and \mathbb{C}_p a lower value indicates a better model. Figure (9) plots these values.

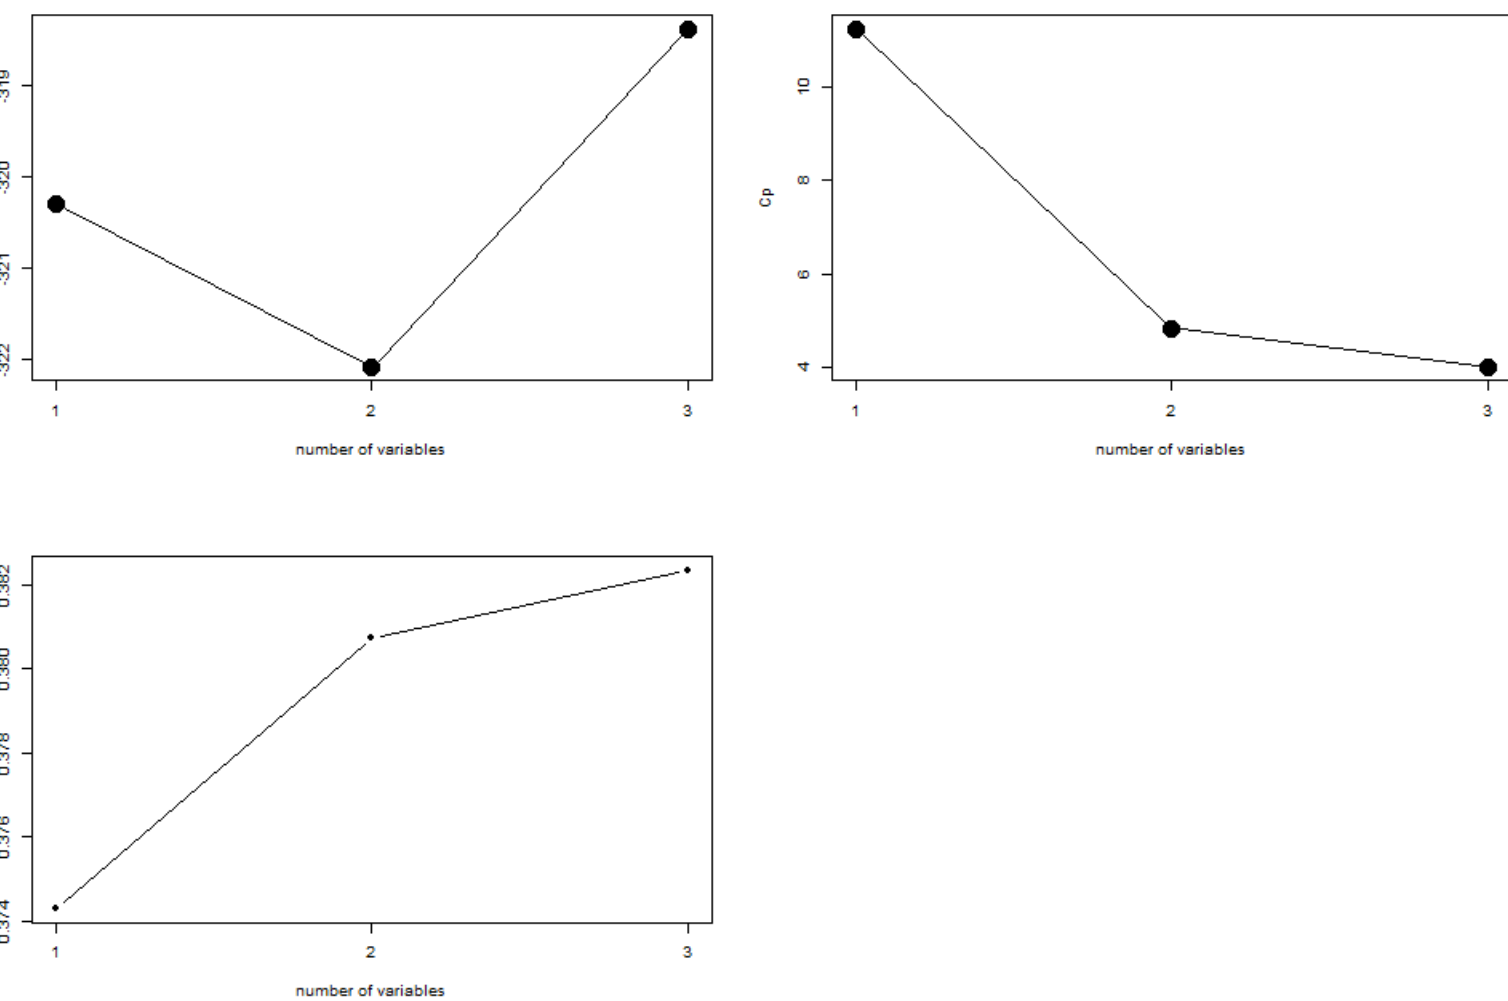


Figure 9: Poly Term Linear Model - Plot of No of Predictor Variables Vs BIC, C_p , R^2

3.3.2 Model Selection - AIC

The most parsimonious model is the model with the lowest AIC value. The *stepAIC()* R call systematically removes one predictor variable at a time till the AIC stops reducing. The following results of the *stepAIC()* call show that the 3 predictor variable model has the lowest AIC of -5892. This is slightly lower than the AIC of -5885.44, for the earlier linear regression model without polynomial terms in Toyota in (3.2.2).

```
Step AIC  GM ~ Ford + poly(Toyota, 2) :
Start:   AIC=-5892
GM ~ Ford + poly(Toyota, 2)
```

	Df	Sum of Sq	RSS	AIC
<none>			0.172	-5892
- poly(Toyota, 2)	2	0.0027	0.175	-5885
- Ford	1	0.0924	0.265	-5590

3.3.3 ANOVA

The following are the ANOVA results of the linear regression model with degree 2 polynomial terms in Toyota. The F-Value for the Poly terms in Toyota is much smaller at $\approx .4\%$. Thus we can confidently reject the null hypothesis and take the model with the Poly term in Toyota as significant.

```
ANOVA for the linear regression model : GM ~ Ford + poly(Toyota, 2) :
Analysis of Variance Table
```

Response: GM

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Ford	1	0.1052	0.1052	430.05	<2e-16 ***
poly(Toyota, 2)	2	0.0027	0.0014	5.62	0.0038 **
Residuals	705	0.1724	0.0002		

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

The following are ANOVA the results comparing the 2 models. Again the F-Value for the Poly terms model is much smaller at $\approx .35\%$. Thus we can confidently reject the null hypothesis and take the linear model with the Polynomial term in Toyota as significant.


```

Comparing Models by ANOVA:
Analysis of Variance Table
Model 1: GM ~ Ford + Toyota
Model 2: GM ~ Ford + poly(Toyota, 2)
  Res.Df  RSS Df Sum of Sq    F Pr(>F)
1     706 0.174
2     705 0.172  1     0.0021 8.58 0.0035 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

3.3.4 Diagnostics - Leverage

The number of high leverage points with this model is 41, ie, there 41 high leverage points with leverage $\geq \frac{2(p+1)}{n}$ where $\frac{(p+1)}{n}$ is the average leverage. This is less than the 59 high leverage points that we had with the earlier model in (3.2.6)

3.3.5 Diagnostics - Studentized Residuals

The plot of the Studentized Residuals for each point Versus the Leverage and Cooks Distance for that point in figure (10) shows much less spread than a similar plot for the earlier model in (3.2.8)

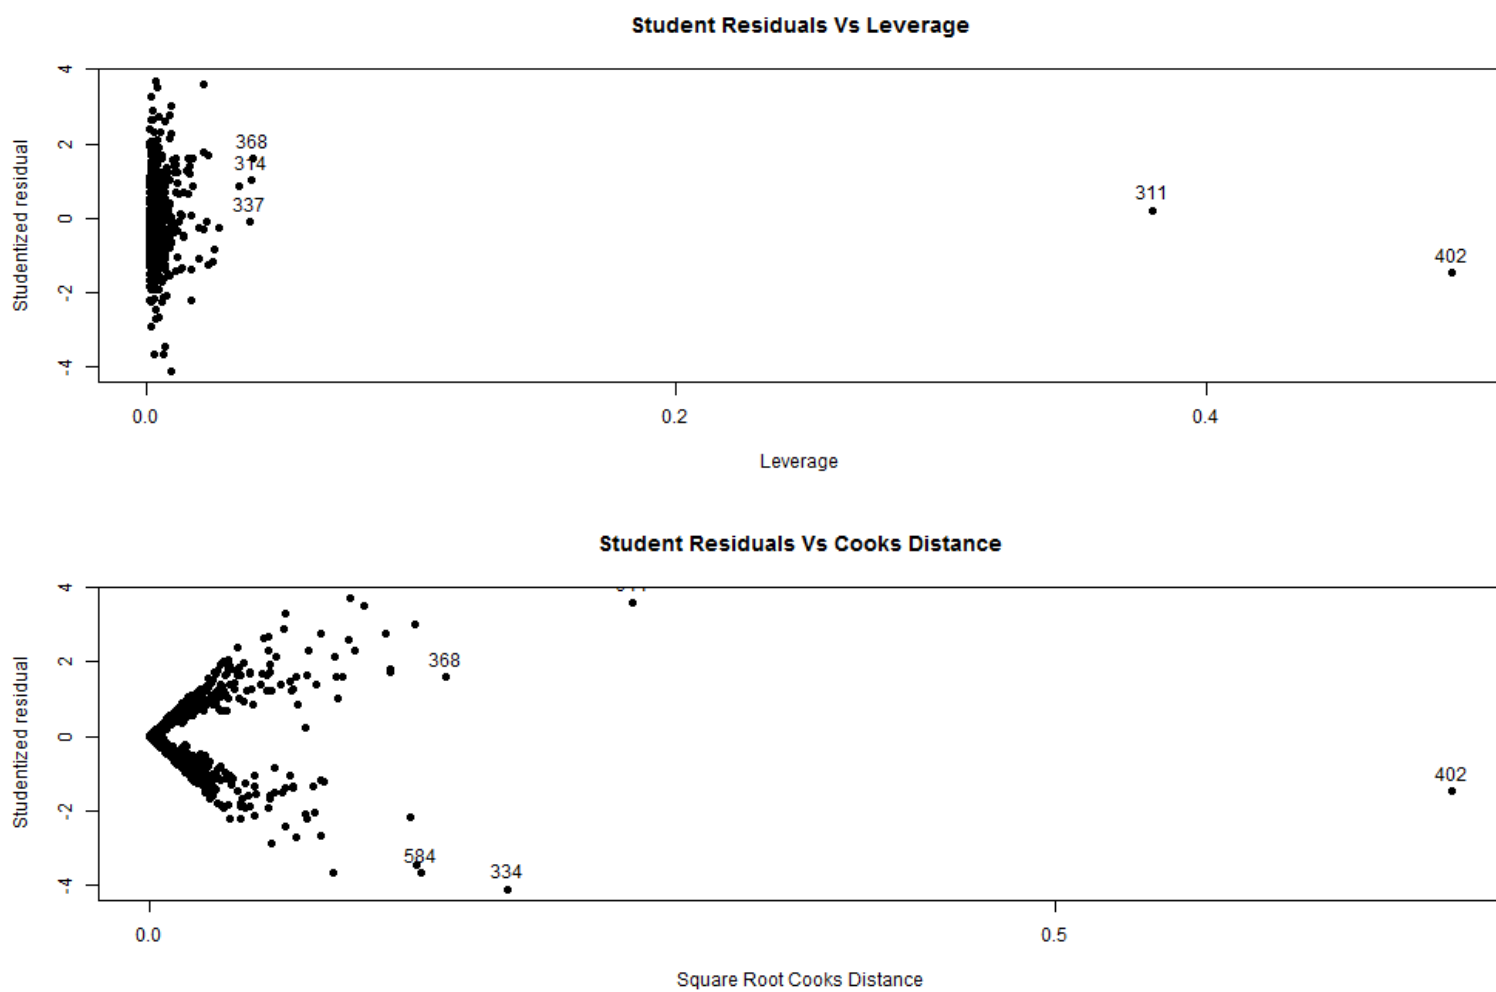


Figure 10: Poly Term Linear Model - Studentized Residuals Versus Leverage and Cooks Distance

3.3.6 Model Assumptions

Figure (11) plots the metrics that enable checking of the model assumptions of normality, linearity and constant variance. The plot of the Studentized Residuals versus Toyota as well as the Studentized Residuals versus Fitted Values shows better linearity with this model than the similar plot for the earlier model in section (3.2.9). This indicates better linearity and better constant variance with this model.

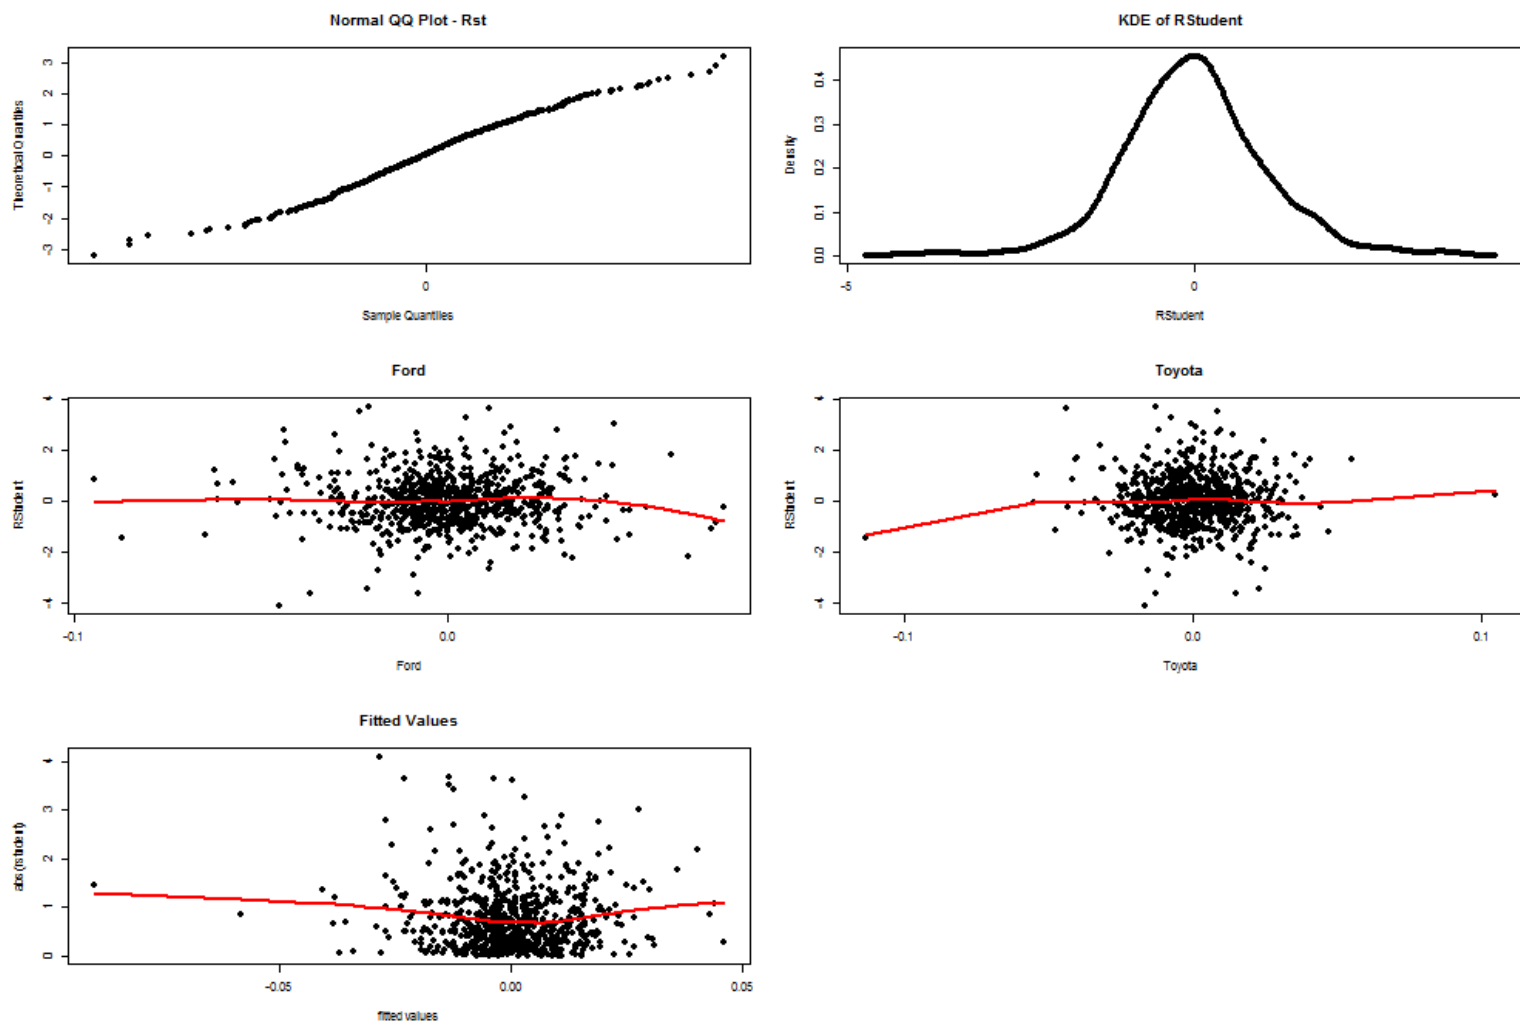


Figure 11: Poly Term Linear Model - Plots to check assumptions of Normality, Linearity and Constant Variance - Plot of Studentized Residuals Vs Predictors, Fitted Values

3.3.7 Compare Model Summary

The following is the summary of the linear regression model degree 2 Polynomial terms in Toyota. As can be seen the P-value is $\approx 0.35\%$ for the 2nd degree polynomial term in Toyota. Thus we can confidently reject the null hypothesis and accept the 2nd degree polynomial term as significant to the model. The $R^2 \approx 0.39$ for this model. This is a slight improvement compared to the $R^2 \approx 0.38$ value for the earlier model in (3.2.10).

Summary for the linear regression model : `GM ~ Ford + poly(Toyota, 2)` :

Call:

```
lm(formula = GM ~ Ford + poly(Toyota, 2))
```

Residuals:

Min	1Q	Median	3Q	Max
-0.06315	-0.00974	-0.00046	0.00866	0.05712

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.000124	0.000587	0.21	0.8324
Ford	0.607437	0.031248	19.44	<2e-16 ***
poly(Toyota, 2)1	0.027149	0.016123	1.68	0.0927 .
poly(Toyota, 2)2	-0.045947	0.015686	-2.93	0.0035 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.016 on 705 degrees of freedom

Multiple R-squared: 0.385, Adjusted R-squared: 0.382

F-statistic: 147 on 3 and 705 DF, p-value: <2e-16

3.4 10 Fold Cross Validation - Evaluate Models

While the analysis so far indicates that the linear regression model with polynomial terms in Toyota is better than the simple linear regression model with only first degree terms in Ford and Toyota, the robust way to validate these models is by performing a 10 fold cross validation on them. In each fold the model is trained on the training set and the results are recorded on the hold out validation set. The Mean of the residual sum of squares across all the validation sets is computed as MSE. This statistic MSE is used to

compared the performance of the models. The following are the MSE results for the two models using 10 fold cross validation. The Figure (12) plots the MSE for each of the models at each fold of cross validation.

```
Results of 10 Fold Cross Valiation: MSE on the Validation Set for Each Model:  
Model : GM ~ Ford + Toyota  
MSE= 0.000251  
Model : GM ~ Ford + poly(Toyota, 2)  
MSE= 0.000248
```




Figure 12: 10 Fold Cross Validation Results - Comparison of Poly(Toyota,2) and Non Poly Term Models - Validation Set MSE for each Fold

4 Conclusion

- As seen from the results of the 10 fold cross validation in section (3.4) the linear regression model with 2nd degree Poly terms in Toyota has a slightly better performance. It has a lower MSE of 0.000248 compared to a MSE of 0.000251 for the linear regression model without any Poly terms. The plot (12) of the MSE at each fold in the cross validation also corroborates that the linear regression model with 2nd degree Poly terms in Toyota has a lower MSE for most of the folds.
- The Poly term in the linear regression model increases the model complexity by adding an additional parameter. So the models need to be compared using a measure of relative model quality like AIC. From the AIC values in section (3.3.2) we see that the linear regression model with degree 2 polynomial terms in Toyota as a AIC of -5892 compared to AIC of -5885.44, for the linear regression model without polynomial terms in Toyota in (3.2.2). Thus the Linear Regression Model with degree 2 polynomial terms in Toyota has a slightly better AIC.
- In section (3.3.7) we see that the Linear Regression Model with degree 2 polynomial terms in Toyota has $R^2 \cong 0.39$ which is a slight improvement compared to the $R^2 \cong 0.38$ value for the linear regression model without polynomial terms.
- We see in section (3.3.4) that the Linear Regression Model with degree 2 polynomial terms in Toyota has fewer high leverage points than the linear regression model without polynomial terms.
- In section (3.3.6) we see that that the Linear Regression Model with degree 2 polynomial terms in Toyota has better linearity and constant variance characteristics than the linear regression model without polynomial terms.
- Thus we conclude that the following Linear Regression Model with degree 2 polynomial terms in Toyota is the better model for the sample data set provided for predicting the log returns for GM , from the log returns of Ford and Toyota as the predictors.

```
GM ~ Ford + poly(Toyota, 2):
llfit$coefficients
      (Intercept)      Ford poly(Toyota, 2)1 poly(Toyota, 2)2
      0.00012      0.60744      0.02715      -0.04595
```


Appendices

The source code for data analysis, for generating the 2 linear regression models and for evaluating the models is in file *hw5_1.R*. This source code is also attached below.

```
#####
##### ISYE6783 #####
##### DL Ajay DSouza #####
##### HW 5 #####
#####

library(leaps)
library(faraway)
library(MASS)
library(car)
library(reshape2)
library(ggplot2)

#-----
# Setup
#
#-----

#cleanup
rm(list = ls())

# set rounding digits globally
options(digits = 6)

set.seed(100)

#setwd(
# "C:/wk/odrive/Amazon Cloud
# Drive/ajays_stuff/georgia_tech_ms/isy_math6783_financial_data_analysis/hw5"
#)

sink("hw5_script_output.txt", append = FALSE, split = TRUE)
sink()
```

```

sink("hw5_script_output.txt", append = TRUE, split = TRUE)
cat('\nHW5\n')
sink()

# Read the ascii data file
data <-
  read.csv(file = "w_logret_3automanu.csv", header = FALSE, sep =
            ",")

# N data size
n <- nrow(data)

# set column names
colnames(data) <- c('Toyota', 'Ford', 'GM')

#attach data
detach(data)
attach(data)

# -----
# Data Analysis
# -----
# 1 scatter plot
par(mfrow = c(1, 1))
scatterplotMatrix( ~ GM + Ford + Toyota, main = "")
dev.copy(png,
          filename = "hw5_g_pairs.png",
          width = 1000,
          height = 700)
dev.off()

# correlation
sink("hw5_script_output.txt", append = TRUE, split = TRUE)
cat(paste("\n\nCorrelation:\n"))
print(cor(data))
sink()

sink("hw5_script_output.txt", append = TRUE, split = TRUE)

```

```

cat(paste("\n\nCorrelation With poly(Toyota,2):\n"))
print( cor(cbind(poly(Toyota,2),Ford,GM)))
sink()

# 2 Histogram, KDE
par(mfrow = c(3, 1), oma = c(0, 0, 2, 0))
den <- density(data$Toyota)
plot(den$x, den$y, xlab = "Toyota", ylab = "Density")

den <- density(data$Ford)
plot(den$x, den$y, xlab = "Ford", ylab = "Density")

den <- density(data$GM)
plot(den$x, den$y, xlab = "GM", ylab = "Density")

#title("Density(KDE) Plot - Daily Log Returns", outer = TRUE)

dev.copy(png,
          filename = "hw5_g_kde.png",
          width = 1000,
          height = 700)
dev.off()

# 3 QQplot with normal distribution
# qqplot with T distribution
tqd <- pt((seq(n) - .5) / n, df = 2.4)

par(mfrow = c(3, 2), oma = c(0, 0, 2, 0))
qqnorm(data$Toyota,
        datax = TRUE,
        ylab = "Toyota",
        main = "QQNorm")
qqplot(data$Toyota, tqd, xlab = "Toyota", main = "QQ-TDist")

qqnorm(data$Ford,
        datax = TRUE,
        ylab = "Ford",
        main = "QQNorm")
qqplot(data$Ford, tqd, xlab = "Ford", main = "QQ-TDist")

```

```

qqnorm(data$GM,
       datax = TRUE,
       ylab = "GM",
       main = "QQNorm")
qqplot(data$GM, tqd, xlab = "GM", main = "QQ-TDist")

#title("QQPlot\n Normal , T-Distribution \n Daily Log Returns", outer =
      TRUE)

dev.copy(png,
       filename = "hw5_g_qq.png",
       width = 1000,
       height = 700)
dev.off()

#-----
# Linear Regression - Model 1 - Analysis
# lm (GM~Ford+Toyota)
#
#-----

for (mdel in c(1, 2)) {
  #anova
  if (mdel == 1) {
    llfit <- lm(data = data, GM ~ Ford + Toyota)
  }
  else if (mdel == 2) {
    lofit <- llfit
    llfit <- lm(GM ~ Ford + poly(Toyota, 2))
  }

  anlfit <- anova(llfit)

  sink("hw5_script_output.txt",
       append = TRUE,
       split = TRUE)
  cat(paste(
    "\n\nANOVA for the linear regression model :",
    llfit$call[2],

```

```

    ":\n"
  ))
  print(anlfit)
  sink()

#summary
smlfit <- summary(llfit)
sink("hw5_script_output.txt",
      append = TRUE,
      split = TRUE)
cat(paste(
  "\n\n\nSummary for the linear regression model :",
  llfit$call[2],
  ":\n"
))
print(smlfit)
sink()


# Anova comparison of two models and F
if (mdel == 1) {
  llfit2 <- lm(data = data, GM ~ Ford)
  acomp <- anova(llfit2, llfit)

} else if (mdel == 2) {
  acomp <- anova(lofit, llfit)
}

sink("hw5_script_output.txt",
      append = TRUE,
      split = TRUE)
cat(paste("\n\n\nComparing Models by ANOVA:\n"))
print(acomp)
sink()


# Model Selection
if (mdel == 1) {
  subsets = regsubsets(data$GM ~ .,

```

```

                                data = as.data.frame(cbind(data$Ford,
                                                                data$Toyota)), nbest =
                                                                1)
} else if (mdel == 2) {
    subsets = regsubsets(data$GM ~ .,
                          data = as.data.frame(cbind(data$Ford,
                                                        poly(data$Toyota, 2))), nbest =
                                                        1)
}

b = summary(subsets)

sink("hw5_script_output.txt",
      append = TRUE,
      split = TRUE)
cat(
  paste(
    "\n\nRegression Subsets - Choosing the Best set of Predictors by
    BIC/Cp:\n"
  )
)
print(b)
sink()


nms <- length(coef(llfit)) - 1

par(
  mfrow = c(2, 2),
  oma = c(0, 0, 2, 0),
  lab = c(2, 5, 3),
  pch = 19
)
plot(
  1:nms,
  b$bic,
  type = "b",
  xlab = "number of variables",
  ylab = "BIC",
  cex = 2.5

```

```

)
plot(
  1:nms,
  b$cp,
  type = "b",
  xlab = "number of variables",
  ylab = "Cp",
  cex = 2.5
)
plot(1:nms,
      b$adjr2,
      type = "b",
      xlab = "number of variables",
      ylab = "adjusted R2")
#title(paste("Model Selection - \nDaily Log Returns ", llfit$call[2]),
#       outer = TRUE)
dev.copy(
  png,
  filename = paste("hw5_m", mdel, "_g_complexity.png", sep = ""),
  width = 1000,
  height = 700
)
dev.off()

#vif for collenarity
vf <- vif(lm(llfit))
sink("hw5_script_output.txt",
      append = TRUE,
      split = TRUE)
cat(paste("\n\nVariance Inflation Factor:\n"))
print(vf)
sink()

#StepAIC

sink("hw5_script_output.txt",
      append = TRUE,
      split = TRUE)

```

```

cat(paste("\n\nStep AIC ", llfit$call[2], ":\n"))
step_lm = stepAIC(llfit)
cat("\n\n")
smlm <- summary(lm(step_lm))
print(smlm)
sink()

#leaps
if (mdel == 1) {
  x1 = as.matrix(cbind(Ford, Toyota))
  names_x1 = c("Ford", "Toyota")
} else if (mdel == 2) {
  x1 = as.matrix(cbind(Ford, poly(Toyota, 2)))
  names_x1 = c("Ford", "Toyota1", "Toyota2")
}

leaps.fit = leaps(
  y = GM,
  x = x1,
  names = names_x1,
  nbest = 1
)
options(digits = 2)

sink("hw5_script_output.txt",
      append = TRUE,
      split = TRUE)
cat(paste("\n\nLeaps - Best Model for each K Predictor by Cp:\n"))
print(cbind(leaps.fit$which, leaps.fit$Cp))
sink()

# Partial Residual Plots
par(
  mfrow = c(2, 2),
  oma = c(0, 0, 2, 0),
  lab = c(2, 5, 3),
  pch = 19
)

```



```

lfit_ford = lm(GM ~ Ford)
lfit_toyota = lm(GM ~ Toyota)
crPlot(
  llfit,
  var = "Ford",
  main = "Partial - Ford",
  smooth = F,
  lty = 1,
  lwd = 2,
  col = "black"
)

if (mdel == 1) {
  crTVar <- 'Toyota'
} else if (mdel == 2) {
  crTVar <- 'poly(Toyota, 2)'
}
crPlot(
  llfit,
  var = crTVar,
  main = paste("Partial-", crTVar),
  smooth = F,
  lty = 1,
  lwd = 2,
  col = "black"
)
plot(Ford, GM, main = "Ford Vs GM")
regLine(lfit_ford,
  col = "red",
  lwd = 2,
  lty = "dashed")
plot(Toyota, GM, main = "Toyota Vs GM")
regLine(lfit_toyota,
  col = "red",
  lwd = 2,
  lty = "dashed")
#title(paste("Partial Residual Plots- \nDaily Log Returns ",
  llfit$call[2]),
#      outer = TRUE)
dev.copy(
  png,

```

```

filename = paste("hw5_m", mdel, "_g_partial_residual.png", sep =
                  ""),
width = 1000,
height = 700
)
dev.off()

# High Leverage Points
hvals <- hatvalues(llfit)
# High leverage points > 2(p+1)/n more than twice the average ( 59 of
  them)
hLevAv <- (length(coef(llfit)) - 1 + 1) / length(llfit$residuals)
hValHAvg <- sum(hvals > 2 * hLevAv)
sIdx <- sort(hvals, decreasing = TRUE, index.return = TRUE)

# Top 5 high leverage points
tp <- 10
sink("hw5_script_output.txt",
      append = TRUE,
      split = TRUE)
cat(paste("\n\n\nTop ", tp, " Data Points with the Highest Leverage:\n"))
print(cbind(data[head(unique(sIdx$ix), tp), ], leverage =
             head(unique(sIdx$x), tp)))

cat(paste(
  "\n\n\nThe Correlation between the ",
  hValHAvg,
  " High Leverage Data Points:\n"
))
print(cor(cbind(data[head(unique(sIdx$ix), hValHAvg), ])))
sink()

#plot
par(
  mfrow = c(2, 1),
  oma = c(0, 0, 2, 0),
  lab = c(2, 5, 3),
  lwd = 1,
  pch = 19
)

```

```

plot(
  hvals,
  ylab = "leverage",
  main = "Leverage of All Points",
  ylim = c(0, .1),
  xlab = "Index"
)
abline(h = 2 * (length(coef(llfit)) - 1 + 1) / length(llfit$residuals),
       col = 'red')

gpts <-
  as.data.frame(cbind(head(unique(sIdx$ix), 3),
    hvals[head(unique(sIdx$ix), 3)]))
colnames(gpts) <- c('Index', 'Leverage')
text(gpts$Index, gpts$Leverage, gpts$Index, pos = 2)

plot(
  which(hvals > 2 * (length(coef(
    llfit
  )) - 1 + 1) / length(llfit$residuals)),
  hvals[hvals > 2 * (length(coef(llfit)) - 1 + 1) /
    length(llfit$residuals)],
  main = paste("Points with Leverage > 2(p+1)/n"),
  ylab = "leverage",
  xlab = "Index"
)
abline(h = 2 * hLevAv, col = 'red')

# Label the top 5 Leverage values
gpts <-
  as.data.frame(cbind(head(unique(sIdx$ix), 5),
    hvals[head(unique(sIdx$ix), 5)]))
colnames(gpts) <- c('Index', 'Leverage')
text(gpts$Index, gpts$Leverage, gpts$Index, pos = 2)

#title(paste("Leverage (Hat Values)- \nDaily Log Returns",
  llfit$call[2]),
  # outer = TRUE)
dev.copy(
  png,

```

```

    filename = paste("hw5_m", mdel, "_g_leverage.png", sep = ""),
    width = 1000,
    height = 700
  )
dev.off()

# RStudent Residuals

# Studentized Residuals
rst <- rstudent(llfit)

# Label the top 5 Leverage values
rlpts <-
  as.data.frame(cbind(head(unique(sIdx$ix), 5),
    hvals[head(unique(sIdx$ix), 5)], rst[head(unique(sIdx$ix), 5)]))
colnames(rlpts) <- c('Index', 'Leverage', 'rst')

# Cook's Distance
sqcdis <- sqrt(cooks.distance(llfit))

# Cook's Distance Plots
sqcdis <- sqrt(cooks.distance(llfit))

# Label the top 5 Leverage values
cIdx <- sort(sqcdis, decreasing = TRUE, index.return = TRUE)

clpts <-
  as.data.frame(cbind(head(unique(cIdx$ix), 5),
    hvals[head(unique(cIdx$ix), 5)], rst[head(unique(cIdx$ix), 5)],
    sqcdis[head(unique(cIdx$ix), 5)]))
colnames(clpts) <- c('Index', 'Leverage', 'rst', 'Cooks')

par(
  mfrow = c(2, 1),
  oma = c(0, 0, 2, 0),
  lab = c(2, 5, 3),
  lwd = 1,
  pch = 19
)

```

```

#plot(rst, ylab = "Studentized residual", main = "Externally Studentized
  Residuals")
#plot(residuals(llfit), ylab = "residual", main = "Raw Residuals")
plot(hvals,
     rst,
     ylab = "Studentized residual",
     xlab = "Leverage",
     main = "Student Residuals Vs Leverage")

text(rlpts$Leverage, rlpts$rst, rlpts$Index, pos = 3)

plot(sqcds,
     rst,
     ylab = "Studentized residual",
     xlab = "Square Root Cooks Distance",
     main = "Student Residuals Vs Cooks Distance")

text(clpts$Cooks, clpts$rst, clpts$Index, pos = 3)

# title(paste("Studentized Residuals- \nDaily Log Returns ",
  llfit$call[2]),
#       outer = TRUE)

dev.copy(
  png,
  filename = paste("hw5_m", mdel, "_g_rstudent_residuals.png", sep =
    ""),
  width = 1000,
  height = 700
)
dev.off()

```

```

# Cooks Distance Plots

```

```

par(
  mfrow = c(2, 1),
  oma = c(0, 0, 2, 0),
  cex.axis = 1,
  cex.lab = 1,
  lwd = 1,
  pch = 19
)

plot(
  sqcdis,
  ylab = ("square root Cook's D"),
  cex = 1,
  main = "Sqrt of Cooks Distance",
  ylim = c(0, .6)
)
text(clpts$Index,
     clpts$Cooks,
     clpts$Index,
     pos = 2,
     cex = 0.7)

halfnorm(
  sqcdis,
  ylab = ("square root Cook's D"),
  cex = 1,
  main = "Half Normal Plots"
)

#title(paste(
#  "Cook's Distance- Daily Log Returns GM~Ford,Toyota",
#  llfit$call[2]
#),
#outer = TRUE)
dev.copy(
  png,
  filename = paste("hw5_m", mdel, "_g-cooks.png", sep = ""),
  width = 1000,
  height = 700
)
dev.off()

```

```

# Check for non constant variance or heteroskedacity
#
# Plot of absolute Residuals versus predictors and response
#

par(
  mfrow = c(3, 2),
  oma = c(0, 0, 2, 0),
  lab = c(2, 5, 3),
  lwd = 1,
  pch = 19
)
arst <- abs(rst)

qqnorm(rst, datax = T, main = "Normal QQ Plot - Rst")

den <- density(rst)
plot(den$x,
     den$y,
     xlab = "RStudent",
     ylab = "Density",
     main = "KDE of RStudent")

plot(Ford,
     rst,
     main = "Ford",
     ylab = "RStudent",
     xlab = "Ford")
fit2 = loess(rst ~ Ford)
ordx2 = order(Ford)
lines(Ford[ordx2], fit2$fitted[ordx2], col = "red", lwd = 2)

plot(Toyota,
     rst,
     main = "Toyota",
     ylab = "RStudent",
     xlab = "Toyota")
fit3 = loess(rst ~ Toyota)

```

```

ordx2 = order(Toyota)
lines(Toyota[ordx2], fit3$fitted[ordx2], col = "red", lwd = 2)

plot(llfit$fitted,
      arst,
      xlab = "fitted values",
      ylab = "abs(rstudent)",
      main = "Fitted Values")
fit4 = loess(arst ~ llfit$fitted)
ord = order(llfit$fitted)
lines(llfit$fitted[ord], fit4$fitted[ord], col = "red", lwd = 2)

# title(
#   paste(
#     "Check for Variance \nResiduals Vs Predictors/Fitted Values \nDaily
#       Log Returns ",
#     llfit$call[2]
#   ),
#   outer = TRUE
# )
dev.copy(
  png,
  filename = paste("hw5_m", mdel, "_g_heteroskedacity.png", sep =
    ""),
  width = 1000,
  height = 700
)
dev.off()

# Plot the final fit
par(
  mfrow = c(2, 2),
  oma = c(0, 0, 2, 0),
  lab = c(2, 5, 3),
  lwd = 1,
  pch = 19
)

```



```

plot(llfit)

dev.copy(
  png,
  filename = paste("hw5_m", mdel, "_g_fit.png", sep = ""),
  width = 1000,
  height = 700
)
dev.off()

}

#-----
#
# Model Comparison
#
# 10 fold cross validation of the chosen models
#
#-----
# Split data into train and hold out test
#

# Number of folds
nFolds <- 10

#Randomly shuffle the array index
smpIdx <- sample(n)

#Create 10 equally size folds
folds <- cut(seq(1, n), breaks = nFolds, labels = FALSE)

# Store the fold results
fResults <- data.frame(
  model = integer(),
  k = integer(),
  trainN = integer(),
  testN = integer(),
  SSR = double(),
  SSE = double(),
  SST = double(),
  MSE = double(),

```

```

MST = double(),
R2 = double(),
R2A = double(),
MSET = double(),
corFitTrain = double(),
corTestFit = double(),
stringsAsFactors = FALSE
)

for (mdel in c(1:2)) {
  #Perform 10 fold cross validation
  for (k in 1:10) {
    testIndices <- smpIdx[which(folds == k, arr.ind = TRUE)]
    ## pick a random train and test rows
    testData <- data[testIndices,]
    trainData <- data[-testIndices,]

    # perform linear regression
    if (mdel == 1) {
      lfit <- lm(data = trainData, GM ~ Ford + Toyota)
    }
    else if (mdel == 2) {
      # perform linear regression
      lfit <- lm(data = trainData, GM ~ Ford + poly(Toyota, 2))
    }

    # residual - mean square error on the fitted data
    yTrainBar <- mean(trainData$GM)
    nTrain <- length(trainData$GM)

    #SSR
    SSR <- sum((predict(lfit) - yTrainBar) ^ 2)

    #SSE and MSE
    SSE <- sum(lfit$residuals ^ 2)
    MSE <- SSE / df.residual(lfit)

    # SST and MST
    SST <- sum((trainData$GM - yTrainBar) ^ 2)
    MST <- SST / (nTrain - 1)
  }
}

```

```

# MSE on the training data
nTest <- length(testData$GM)
testPred <- predict.lm(lfit, testData[, c('Ford', 'Toyota')])
MSET <- sum((testPred - testData$GM) ^ 2) / nTest

# Save the results
fResults[nrow(fResults) + 1,] <- c(
  mdel,
  k,
  nTrain,
  nTest,
  SSR,
  SSE,
  SST,
  MSE,
  MST,
  SSR / SST,
  1 - (MSE / MST),
  MSET,
  cor(lfit$fitted.values, trainData$GM) **
    2,
  cor(testPred, testData$GM) ** 2
)
}
}

CV_MSE_1 <-
  sum(fResults[fResults$model == 1, ]$MSET * fResults[fResults$model == 1,
    ]$testN) /
  n
CV_MSE_2 <-
  sum(fResults[fResults$model == 2, ]$MSET * fResults[fResults$model == 2,
    ]$testN) /
  n

sink("hw5_script_output.txt", append = TRUE, split = TRUE)
cat(
  paste(

```

```

      "\n\n\n\nResults of 10 Fold Cross Valiation: MSE on the Validation Set
      for Each Model:\n"
    )
  )
  cat(paste("\nModel :", lofit$call[2]))
  cat(paste("\nMSE=", round(CV_MSE_1, 6)))
  cat(paste("\n"))
  cat(paste("\nModel :", llfit$call[2]))
  cat(paste("\nMSE=", round(CV_MSE_2, 6)))
  cat(paste("\n"))
  sink()

fResults$model <- as.factor(fResults$model)

#melt with method name as pivot
molten <-
  melt(
    fResults[, c('k', 'model', 'MSET')],
    id.vars = c('model', 'k'),
    measure.vars = c('MSET'),
    variable_name = 'series'
  )

klabs <- c(paste(unlist(seq(1:10))))

par(
  mfrow = c(1, 1),
  oma = c(0, 0, 2, 0),
  lab = c(2, 5, 3),
  lwd = 1,
  pch = 19
)

ggplot(molten, aes(x = k, y = value, colour = model)) +
  geom_line() +
  scale_x_continuous(breaks = 1:10, labels = klabs) +
  xlab("K-Fold") +
  ylab("MSE") +
  scale_colour_discrete(
    name = "Model",

```

```

    breaks = c("1", "2"),
    labels = c(lofit$call[2], llfit$call[2])
) +
scale_shape_discrete(
  name = "Model",
  breaks = c("1", "2"),
  labels = c(lofit$call[2], llfit$call[2])
) +
ggtitle(paste("10 Fold Validation, MSE on Validation Set")) +
theme_bw() +
theme(plot.title = element_text(hjust = 0.5),
      legend.position = "bottom")

dev.copy(png,
         filename = "hw5_cv_plot.png",
         width = 1000,
         height = 700)
dev.off()

#-----
# END
#-----

```

References

- [1] Ruppert David, Matteson S. David, *Statistics and Data Analysis for Financial Engineering*, Springer-Verlag New York Edition 2, 2015.